



Prof. Jose F. Rodrigues Jr.  
University of Sao Paulo, Brazil

Edited and extended from the slides publicly distributed by:  
A. Im, G. Cai, H. Tunc, J. Stevens, Y. Barve, S. Hei  
Vanderbilt University

<http://conteudo.icmc.usp.br/pessoas/junio/Site/index.htm>



# Content

- ▶ Part 1: Intro and CRUD
  - ▶ 1: Introduction & Basics
  - ▶ 2: CRUD

# History

- ▶ mongoDB = “Humongous DB”
  - ▶ Open-source
  - ▶ Document-based
  - ▶ “High performance, high availability”
  - ▶ Automatic scaling

-[blog.mongodb.org/post/475279604/on-distributed-consistency-part-1](http://blog.mongodb.org/post/475279604/on-distributed-consistency-part-1)  
-[mongodb.org/manual](http://mongodb.org/manual)

# Motivations

- Problems with SQL
  - Rigid schema
  - Not easily scalable (designed for 90's technology or worse)
  - Requires unintuitive joins (despite its claims, Mongo does not do any better due to physical constraints)
- Perks of mongoDB
  - Easy interface with common languages (Java, Javascript, PHP, etc.)
  - Keeps essential features of RDBMS's while learning from key-value noSQL systems

[http://www.slideshare.net/spf13/mongodb-9794741?v=qp1&b=&from\\_search=13](http://www.slideshare.net/spf13/mongodb-9794741?v=qp1&b=&from_search=13)

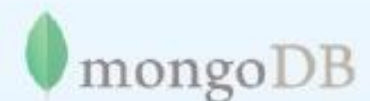
# Company Using MongoDB



“MongoDB powers Under Armour’s online store, and was chosen for its dynamic schema, ability to scale horizontally and perform multi-data center replication.”

<http://www.mongodb.org/about/production-deployments/>

# In Good Company



-Steve Francia, [http://www.slideshare.net/spf13/mongodb-9794741?v=qp1&b=&from\\_search=13](http://www.slideshare.net/spf13/mongodb-9794741?v=qp1&b=&from_search=13)

# Data Model

- Document-Based (max 16 MB each entry)
- Documents are in BSON format, consisting of field-value pairs
- Each document stored in a collection
- Collections
  - Like tables of relational db's.
  - Documents do not have to have uniform structure

-docs.mongodb.org/manual/

# JSON

- “JavaScript Object Notation”
- Easy for humans to write/read, easy for computers to parse/generate
- Objects can be nested
- Built on
  - name/value pairs
  - ordered list of values

<http://json.org/>



# BSON

- “Binary JSON”
- Binary-encoded serialization of JSON-like docs
- Goals
  - Lightweight
  - Traversable
  - Efficient (decoding and encoding)

<http://bsonspec.org/>

# BSON Example

id	user_name	email	age	city
1	Mark Hanks	mark@abc.com	25	Los Angeles
2	Richard Peter	richard@abc.com	31	Dallas



```
{
  "_id": ObjectId("5146bb52d8524270060001f3"),
  "age": 25,
  "city": "Los Angeles",
  "email": "mark@abc.com",
  "user_name": "Mark Hanks"
}
{
  "_id": ObjectId("5146bb52d8524270060001f2"),
  "age": 31,
  "city": "Dallas",
  "email": "richard@abc.com",
  "user_name": "Richard Peter"
}
```

# BSON Example

```
{
  "_id" : "37010"
  "city" : "ADAMS",
  "pop" : 2660,
  "state" : "TN",
  "congressmen:" : ["John", "Willian", "Adolf"]
  "mayor" : {
    name: "John Smith"
    address: "13 Scenic Way"
  }
}
```

➔ Embedding and arrays, more similar to what we have in all-purpose programming languages

# BSON Types

Type	Number
Double	1
String	2
Object	3
Array	4
Binary data	5
Object id	7
Boolean	8
Date	9
Null	10
Regular Expression	11
JavaScript	13
Symbol	14
JavaScript (with scope)	15
32-bit integer	16
Timestamp	17
64-bit integer	18
Min key	255
Max key	127

<http://docs.mongodb.org/manual/reference/bson-types/>

<https://docs.mongodb.com/manual/reference/operator/query/type/>

<http://conteudo.icmc.usp.br/pessoas/junio/Site/index.htm>

# The `_id` Field

- By default, each document contains an `_id` field. This field has a number of special characteristics:
  - Primary key for collection.
  - Value is unique, immutable, and may be any non-array type.
  - Default data type is `ObjectId`, which is “small, likely unique, fast to generate, and ordered.” Sorting on an `ObjectId` value is roughly equivalent to sorting on creation time.

<http://docs.mongodb.org/manual/reference/bson-types/>

# The \_id Field

- Using the default \_id:

```
db.collection.insert({city: "New York", state:"NY", pop:"5M"})
```

- Using your own \_id:

```
db.collection.insert({_id: 10, city: "New York", state:"NY", pop:"5M"})
```

- Using your own composite \_id:

```
db.collection.insert({_id: {city: "New York", state:"NY"}, pop:"5M"})
```

The \_id itself is a document.

<http://docs.mongodb.org/manual/reference/bson-types/>

# mongoDB vs. SQL

mongoDB	SQL
Document	Tuple
Collection	Table/View
PK: _id Field	PK: Any Attribute(s)
Uniformity not Required	Uniform Relation Schema
Index	Index
Embedded Structure	Joins
Shard	Partition
CRUD	DML

## 2. CRUD

Create, Read, Update, Delete



# Getting Started with mongoDB

To install mongoDB, go to this link and click on the appropriate OS and architecture:  
<http://www.mongodb.org/downloads>

First, extract the files (preferably to the C drive).

Finally, create a data directory on C:\ for mongoDB to use

i.e. “md data” followed by “md data\db”

<http://docs.mongodb.org/manual/tutorial/install-mongodb-on-windows/>

# Install

- Unzip to any directory
- Find executable “mongod”
- Default connection at `localhost:27017`
- Run

```
mongod --dbpath .
```

- Or just

```
mongod
```

For the default dir (`/var/lib/mongodb/` or `c:\data\db`)

- Or run from bin dir, and have data anywhere else

```
mongod --dbpath <any dir path>
```

- Visual interface

```
https://www.mongodb.com/products/compass
```

# CRUD: Using the Shell

➔ To establish a connection to the server, open another command prompt window and go to the same directory, entering “mongo.exe”

To check which db you're using

```
db
```

Show all databases

```
show dbs
```

Switch db's/make a new one

```
use <name>
```

See what collections exist

```
show collections
```

Create collection

```
db.createCollection("<name>")
```

Note: db's are not actually created until you insert data!

# CRUD - summary

## SQL to Aggregation Mapping

SQL Terms, Functions, and Concepts	MongoDB Aggregation Operators
WHERE	<code>\$match</code>
GROUP BY	<code>\$group</code>
HAVING	<code>\$match</code>
SELECT	<code>\$project</code>
ORDER BY	<code>\$sort</code>
LIMIT	<code>\$limit</code>
SUM()	<code>\$sum</code>
COUNT()	<code>\$sum</code>

### Mapping Chart:

<http://docs.mongodb.org/manual/reference/sql-aggregation-comparison/>

# CRUD: Using the Shell (cont.)

To insert documents into a collection/make a new collection:

```
db.<collection>.insert(<document>)
```

<=>

```
INSERT INTO <table>  
VALUES(<attributevalues>);
```

# CRUD: Inserting Data

Insert one document

```
db.<collection>.insert({<field>:<value>})
```

Inserting a document with a field name new to the collection is inherently supported by the BSON model.

To insert multiple documents, use an array.

# CRUD: Querying

- Get all docs: `db.<collection>.find()`
  - Returns a cursor, which is iterated to display first 20 results.
  - Add “`.limit(<number>)`” to limit results  
`db.<collection>.find().limit(2)`
  - `SELECT * FROM <table>;`
- Get one doc: `db.<collection>.findOne()`, the first in the disk physical order, usually the first inserted

# CRUD: Querying

To match a specific value:

```
db.<collection>.find({<field>:<value>})
```

“AND”:

```
db.<collection>.find({<field1>:<value1>,  
                     <field2>:<value2>  
                     })
```

**SELECT \***

**FROM <table>**

**WHERE <field1> = <value1> AND <field2> = <value2>;**



# CRUD: Querying

OR

```
db.<collection>.find({ $or: [  
  {<field>:<value1>},  
  {<field>:<value2>}  
] })
```

```
SELECT *  
FROM <table>  
WHERE <field> = <value1> OR <field> = <value2>;
```

Checking for multiple values of a set:

```
db.<collection>.find({<field>: {$in [<value>, <value>]}})
```

```
SELECT *  
FROM <table>  
WHERE <field> IN (<value>,<value>);
```

# CRUD: Querying

Including/excluding document fields

```
db.<collection>.find({ }, {<field1>: 1})
```

```
SELECT field1
```

```
FROM <table>;
```

↓  
0 false  
>0 true

```
db.<collection>.find({<field1>:<value>}, {<field1>: 1})
```

```
SELECT field1
```

```
FROM <table>
```

```
WHERE <field1> = <value>;
```

# CRUD: Querying

Including/excluding document fields

```
db.<collection>.find({<field1>:<value>}, {<field2>: 0})
```

**SELECT <all fields but not field2>**

**FROM <table>**

**WHERE <field1> = <value>;**

- notice that find() takes two parameters

# CRUD: Updating

```
db.<collection>.update(  
  {<field1>:<value1>},    //all docs in which field = value  
  {$set: {<field2>:<value2>}},    //set field to value  
  {multi:true} )           //update multiple docs
```

```
UPDATE <table>  
SET <field2> = <value2>  
WHERE <field1> = <value1>;
```

# CRUD: Updating

To remove a field

```
db.<collection>.update({<field>:<value>},  
    { $unset: { <field>: 1}})
```

```
ALTER TABLE DROP COLUMN <field>  
“WHERE field = value”
```

# CRUD: Removal

Remove all records where field = value

```
db.<collection>.remove({<field>:<value>})
```

```
DELETE FROM <table>  
WHERE <field> = <value>;
```

As above, but only remove first document

```
db.<collection>.remove({<field>:<value>}, true)
```

# CRUD: Isolation

- By default, all writes are atomic **only** on the level of a single document.
- This means that writes over multiple documents of the same collection can be interleaved with other operations.
- You can isolate writes on an **entire** collection by adding “**\$isolated:1**” in the query area:

```
db.foo.update(  
  { status : "A" , $isolated : 1 },  
  { $inc : { count : 1 } },  
  { multi: true }  
)
```

search criterium

--increments by 1 the field *count* of every document  
--with *status A* in the collection foo

➔ In this example, the **\$isolated :1** clause makes other clients wait to read and to write the collection until the command is completed

# Access control included

Authentication mode must be set during start up:  
`mongod --auth`

Then, users must be created:

```
use admin /*as administrator*/ password
db.createUser(
  { user: "reportsUser", pwd: "12345678",
    roles: [
      { role: "read", db: "reporting" },
      { role: "read", db: "products" },
      { role: "readWrite", db: "accounts" }
    ]
  }
)
```



## For More Information

Resource	Location
MongoDB Downloads	<a href="http://mongodb.com/download">mongodb.com/download</a>
Free Online Training	<a href="http://education.mongodb.com">education.mongodb.com</a>
Webinars and Events	<a href="http://mongodb.com/events">mongodb.com/events</a>
White Papers	<a href="http://mongodb.com/white-papers">mongodb.com/white-papers</a>
Case Studies	<a href="http://mongodb.com/customers">mongodb.com/customers</a>
Presentations	<a href="http://mongodb.com/presentations">mongodb.com/presentations</a>
Documentation	<a href="http://docs.mongodb.org">docs.mongodb.org</a>
Additional Info	<a href="mailto:info@mongodb.com">info@mongodb.com</a>

**MongoDB: The Definitive Guide,  
By Kristina Chodorow and Mike Dirolf**

**Published: 9/24/2010**

**Pages: 216**

**Language: English**

**Publisher: O'Reilly Media, CA**

